

Parallel Cartographic Modeling Language : A 30 Minute Introduction to Spatial Data Processing in Python

Eric Shook

High-Performance Computing and GIS Lab

Department of Geography

Kent State University

eshook@kent.edu



What we will cover

Review Cartographic Modeling

Introduce Parallel Cartographic Modeling Language

Create our own Local Sum Operation

You will join the ranks of awesomeness by becoming a parallel programmer

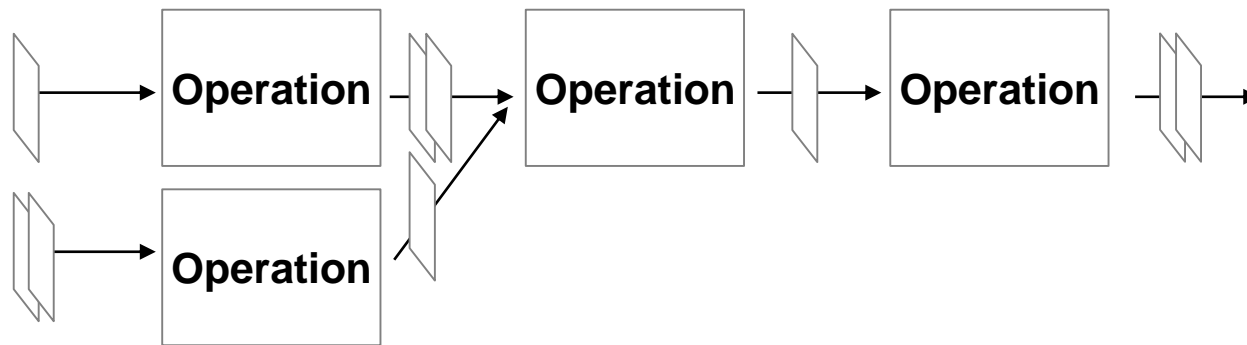
Explore spatial data-processing using PCML



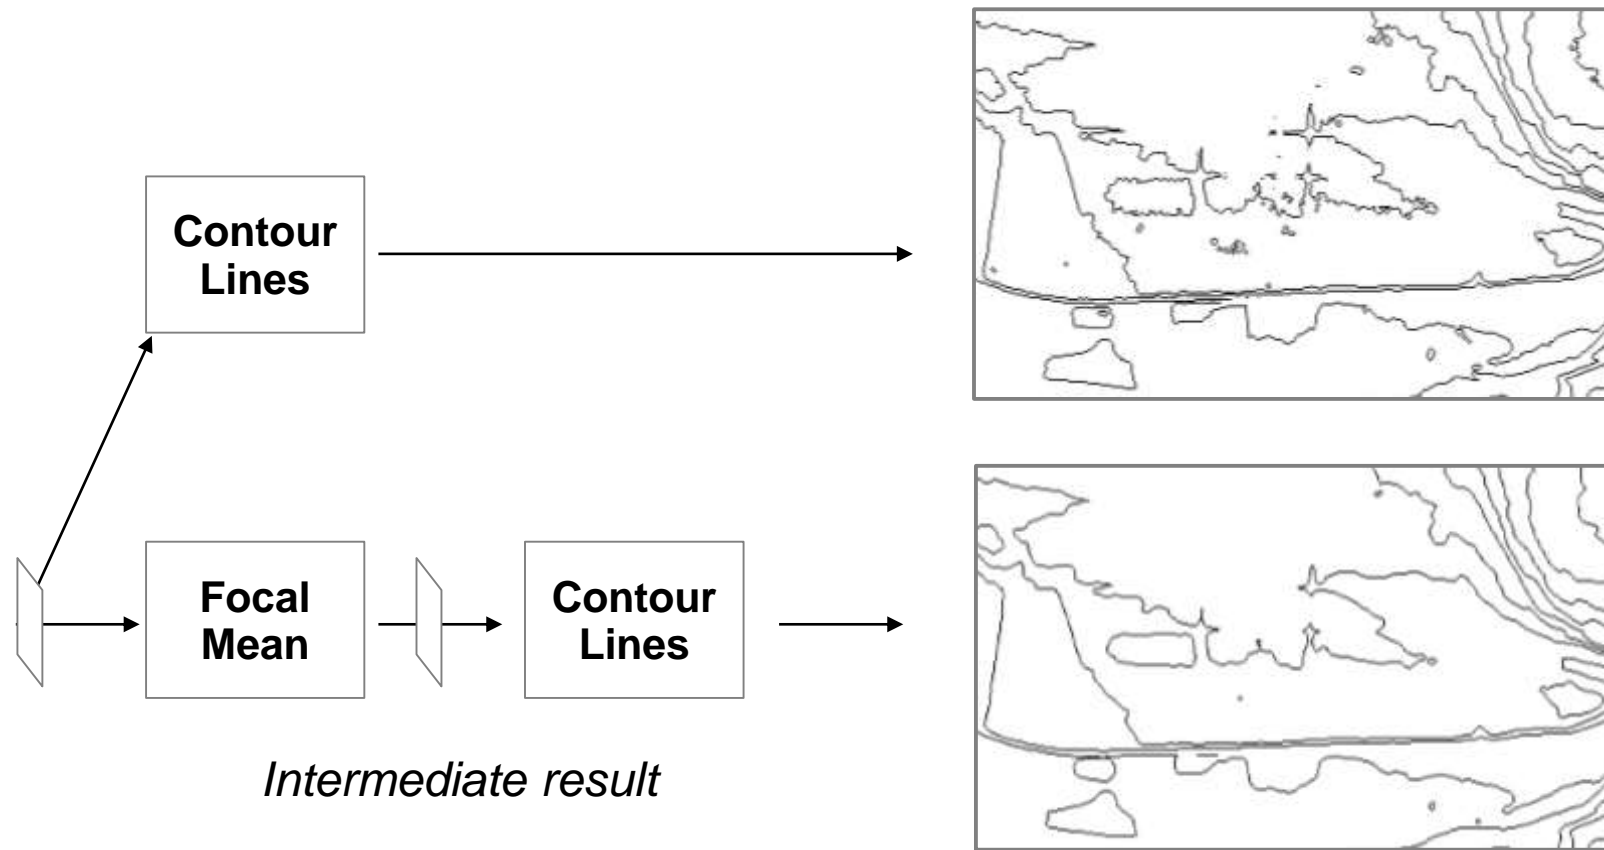
Cartographic Modeling

"Cartographic Modeling is a general methodology for the analysis and synthesis of geographical data"
(Tomlin 1990)

- Clearly and consistently decomposes spatial data processing into elementary components
- Often expressed as a flowchart




Cartographic Modeling: Simple Example



Many operations exist from buffers to calculating hill slopes to reclassifications

Cartographic Modeling: Simple Example



Rather than just connecting operations together

We will create our own operation and join the data-intensive computing party (yay!)

myLocalSum – What we are going to write

$$1 + 2 = 3$$

layera

1	2
3	4

+

layerb

2	4
1	6

=

myresult

3	6
4	10

Our own Local Sum Operation


Parallel Cartographic Modeling Language (PCML)
was created based on 3 design goals:

1. Usability *(Easy to use)*
2. Programmability *(Easy to develop)*
3. Scalability *(Able to handle big data)*

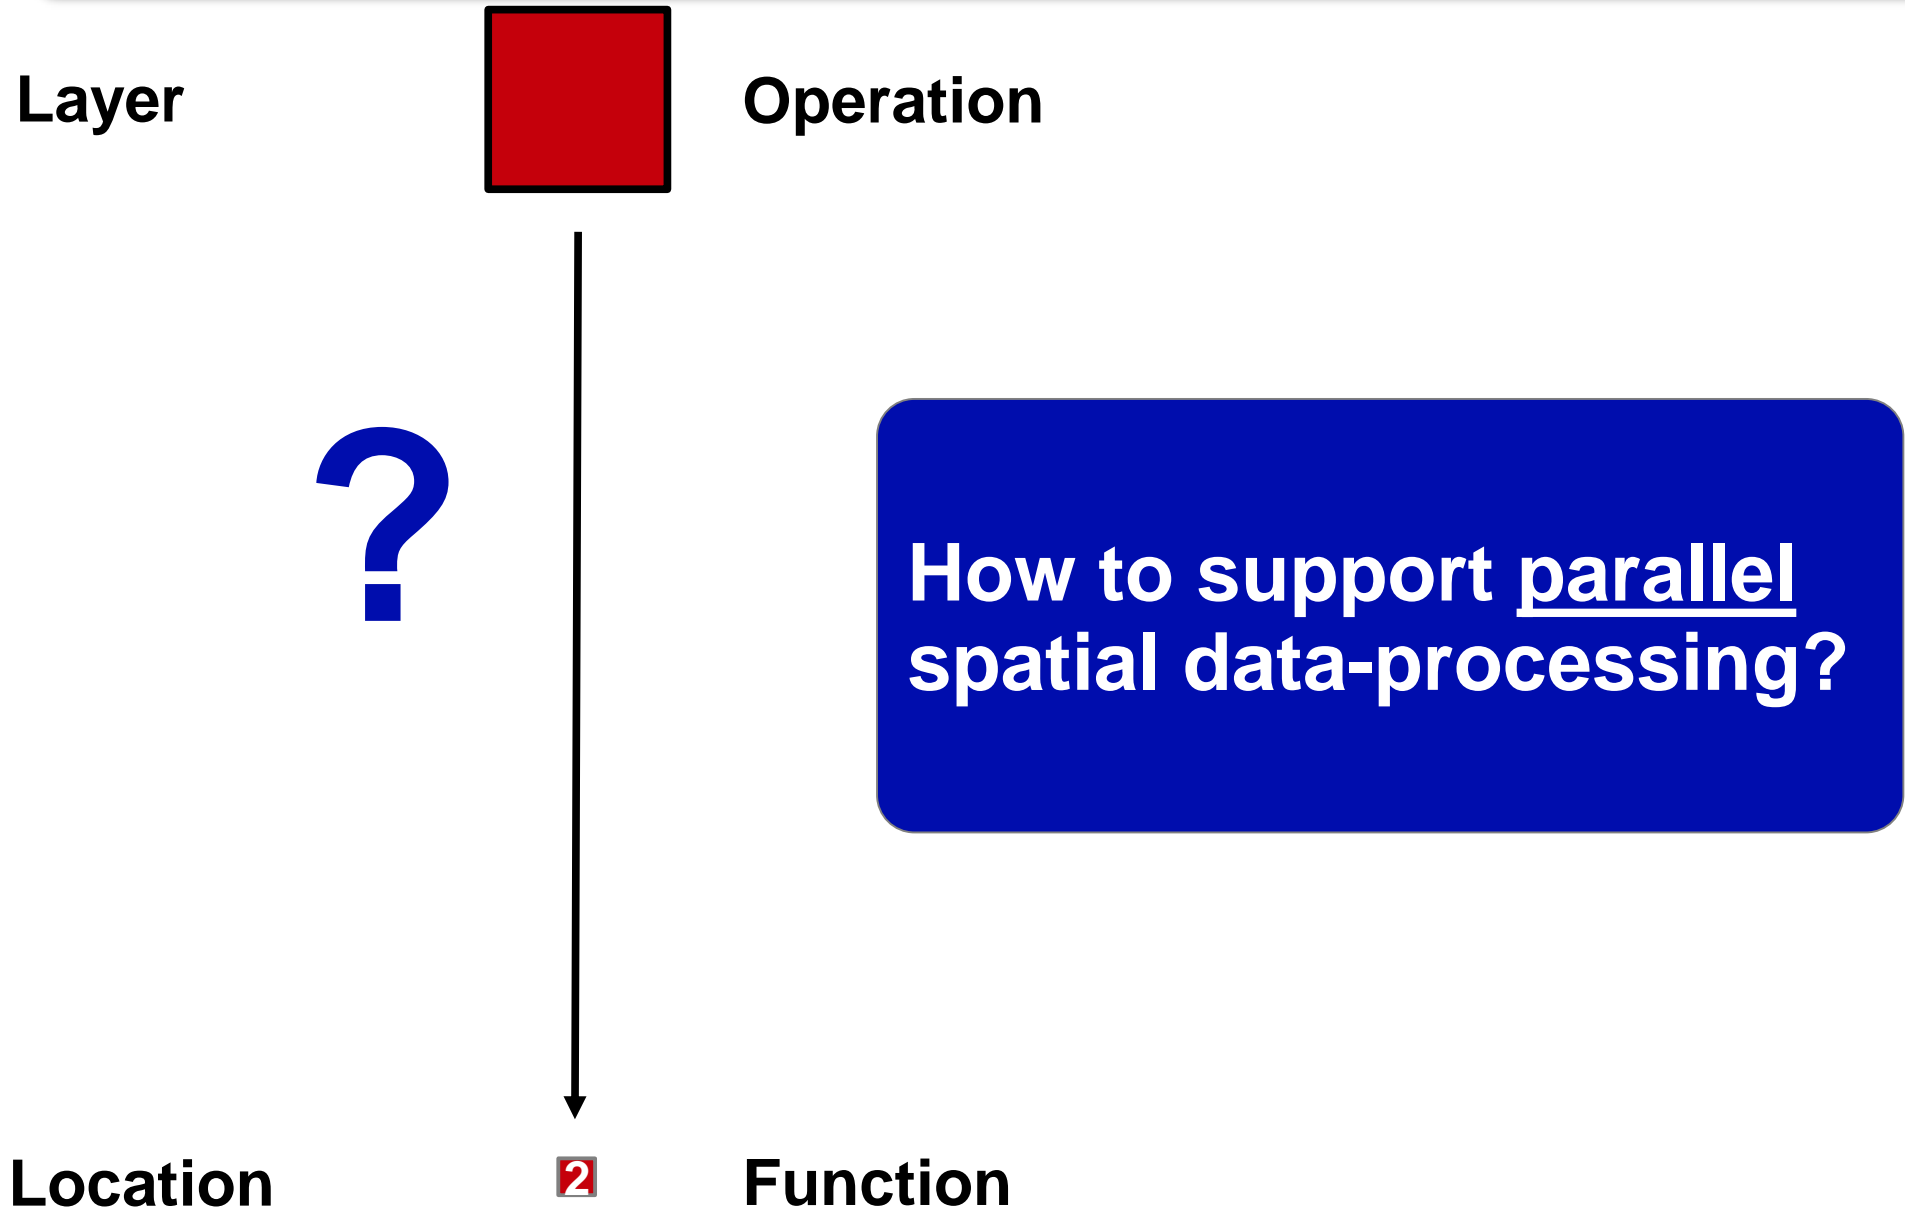
PCML supports automatic parallelization



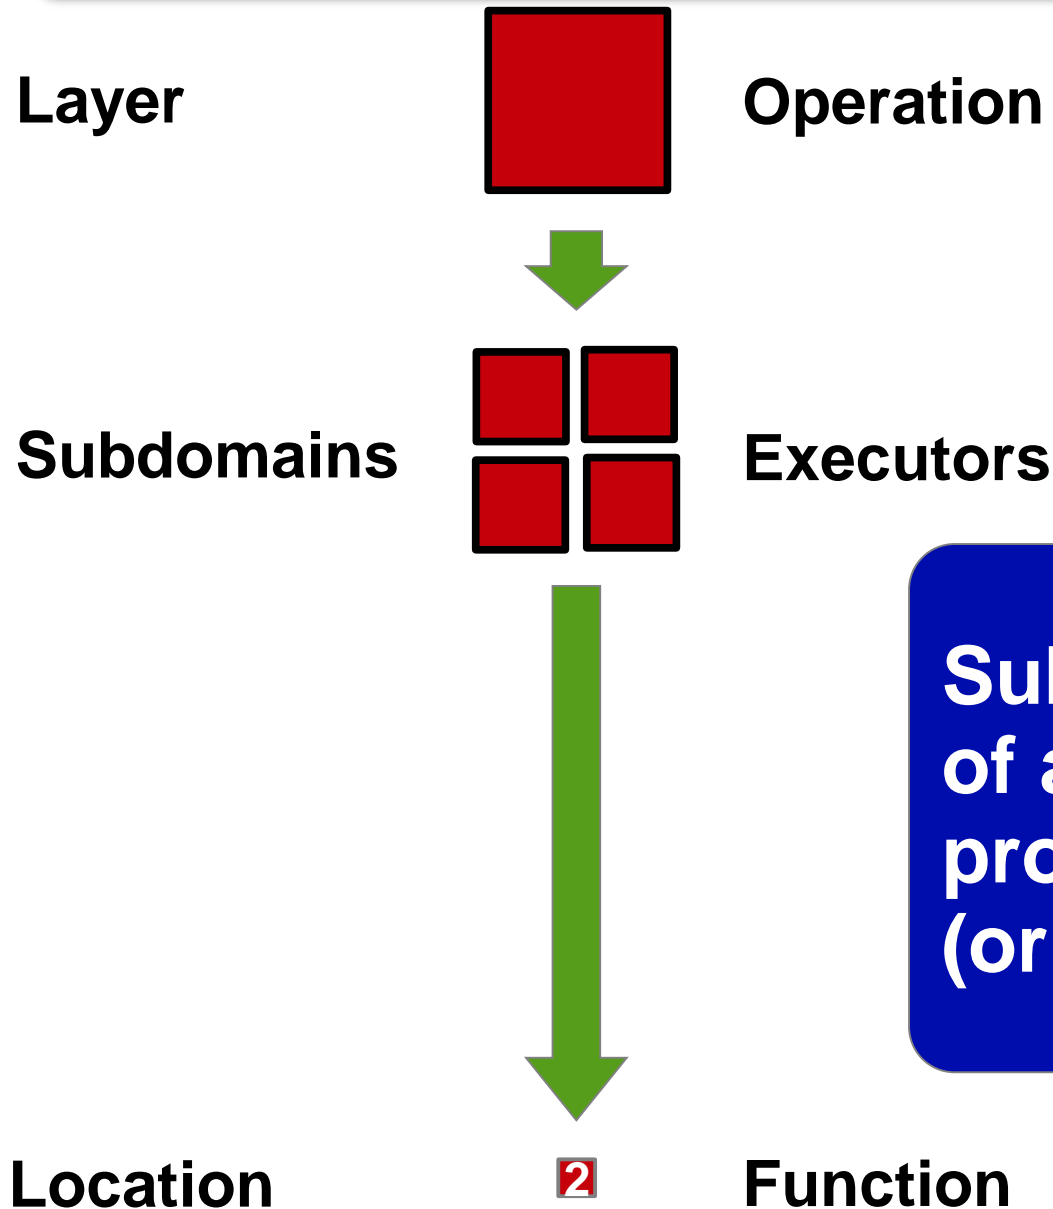
Cartographic Modeling



Cartographic Modeling

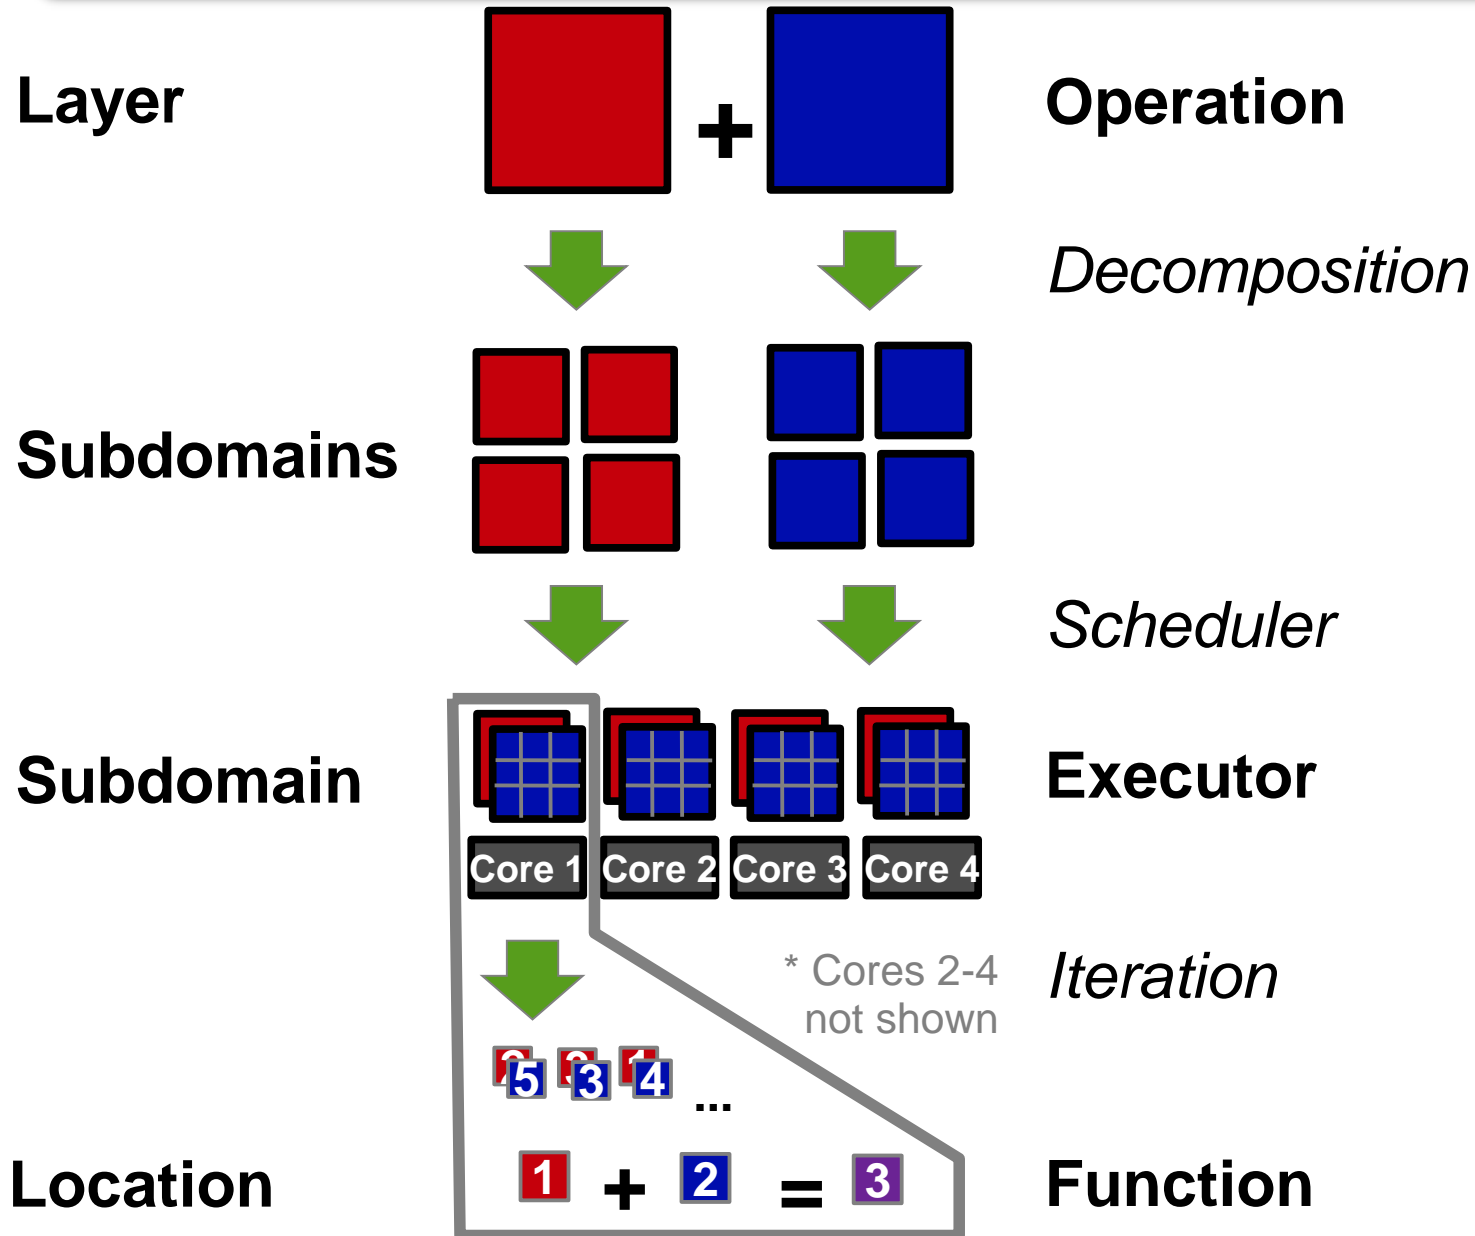


Subdomains

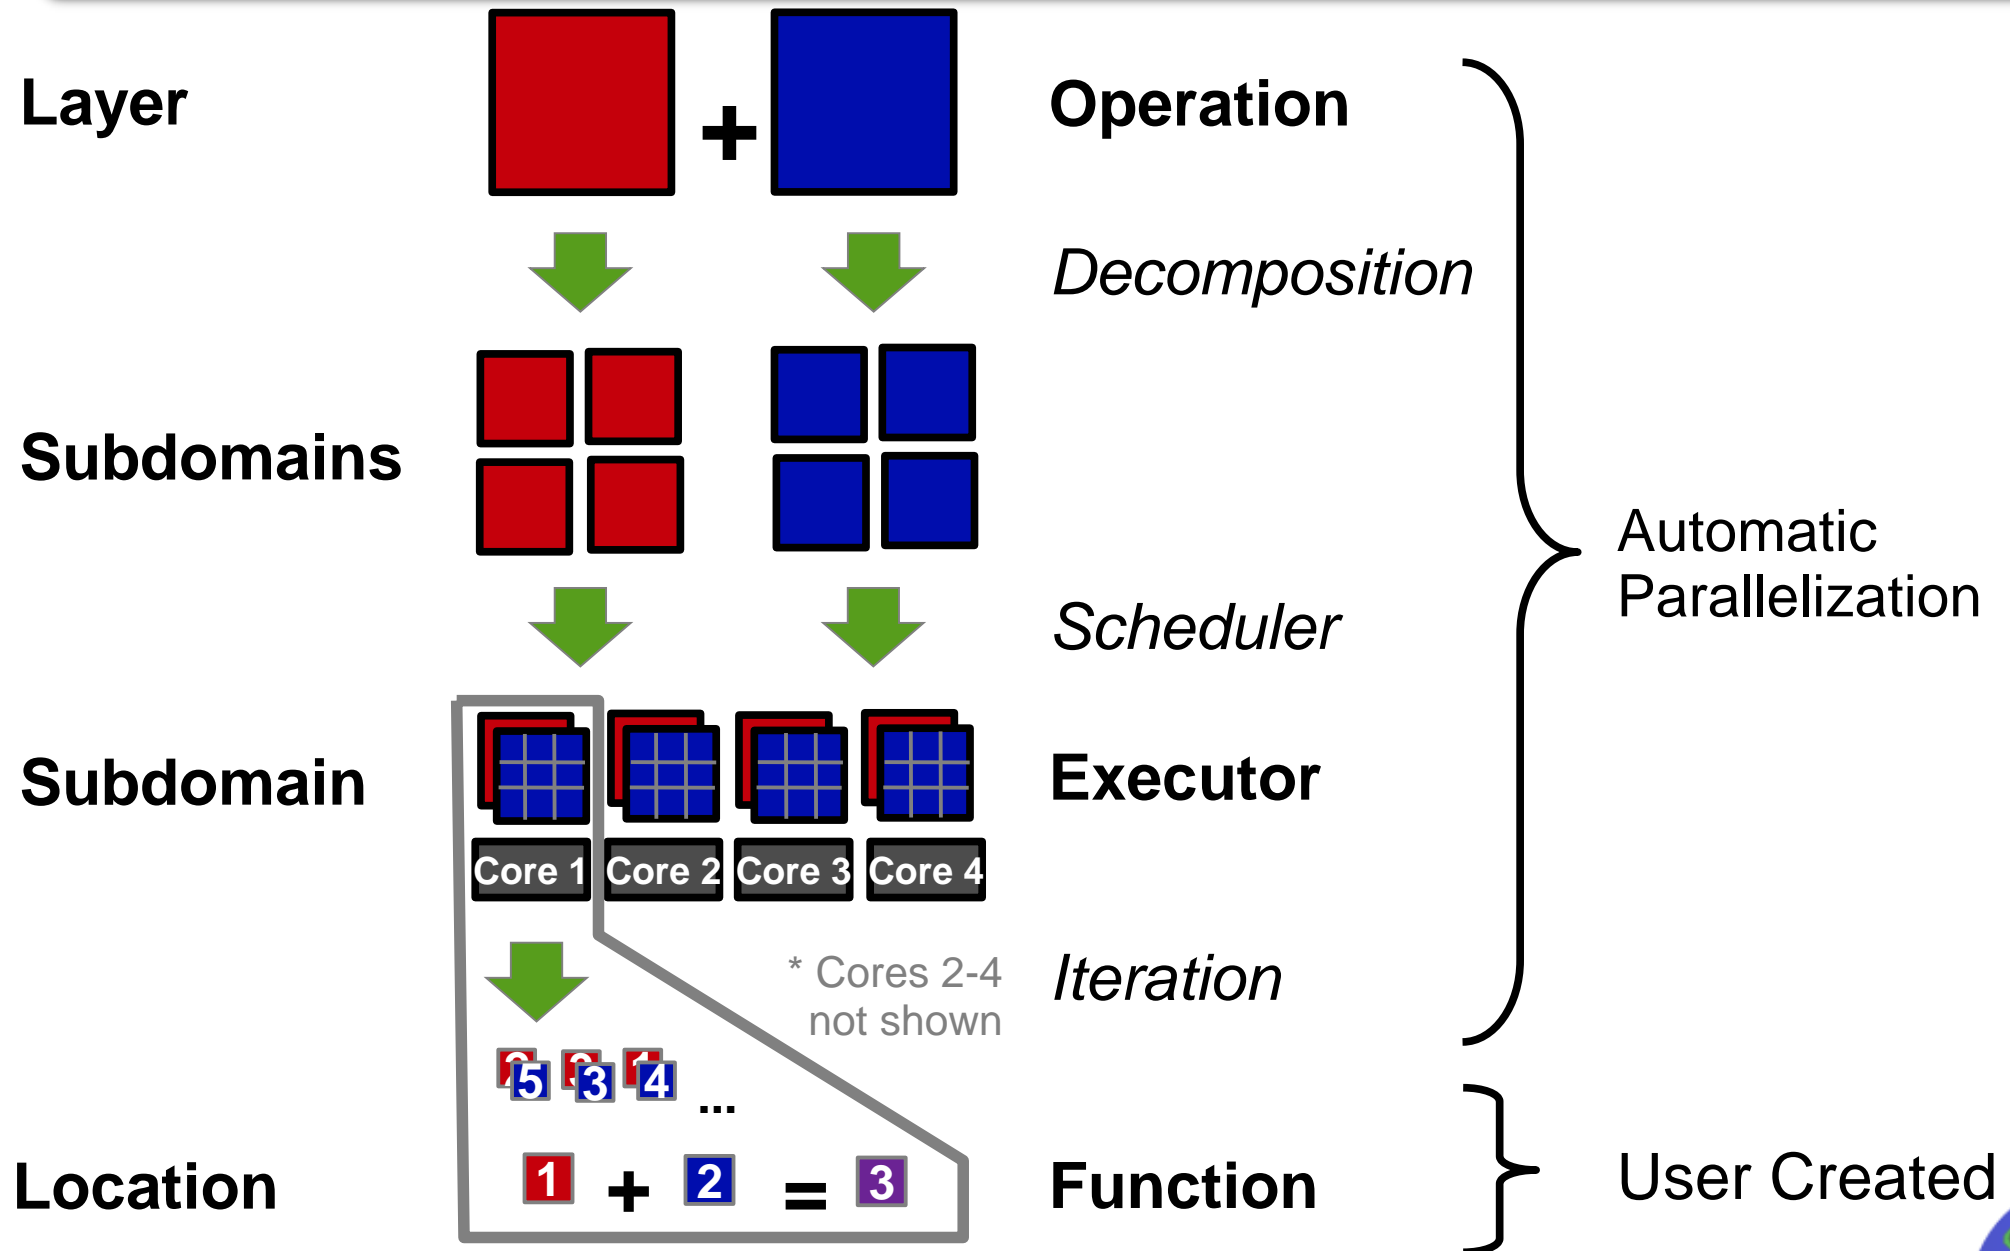


Subdomains are portions of a layer that can be processed simultaneously (or in parallel) in PCML

PCML Conceptual Design



PCML Conceptual Design




Step 0 – Download PCML

<https://github.com/hpcgislab/pcml>

- Click "Download ZIP"
- Unzip on your computer
- Open:

demo_30min.py



Okay let's try it!

30 Minute Demo Overview

(demo_30min.py)

Read 2 data layers (4 x 6 cell raster files)

Print the values out to the screen

Add them together

You will write a myLocalSum Operation

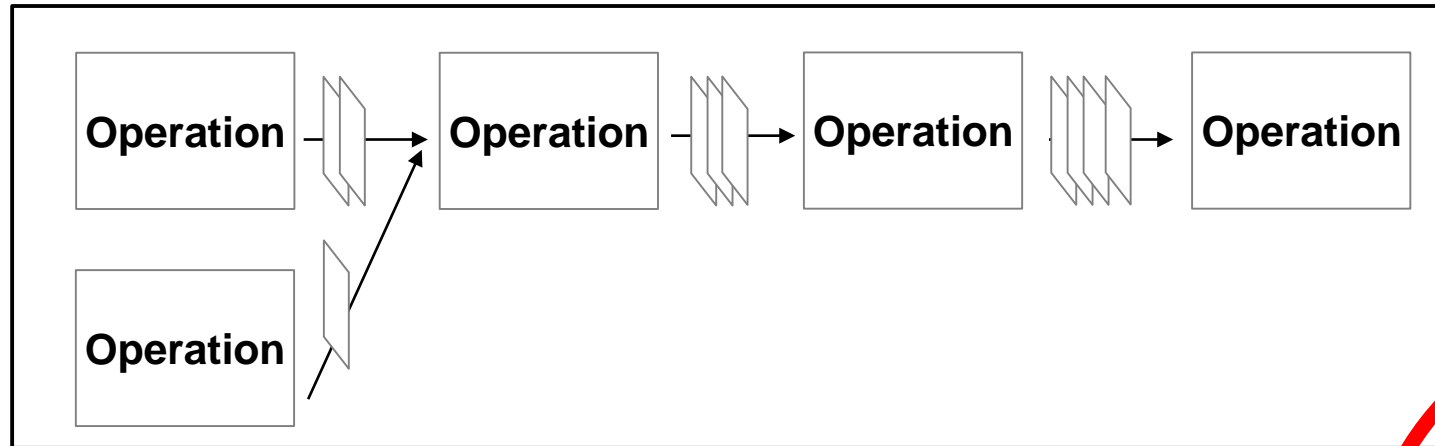
You will use your operation to sum them

You will compare the results

You will then be an awesome
PCML programmer!


Review: Cartographic Modeling

Procedure is a sequence of Operations



Layer

Location



Operation



Function

Classification:



Local
Focal
Zonal
Global

Local, Focal, Zonal, Global Operations



Local Operation





Focal Operation



Zonal Operation



Global Operation



1a – Classify your Operation

```
#####  
#  
# 1. In this section, you will create your own myLocalSum operation  
#  
#####
```

@localoperation

Decorator

Decorators are easy to use. They allow us extend the functionality of Python. We will extend the functionality of a *local operation*, which is a built-in component of PCML, which provides automatic parallelization.

1b – Define your Function

```
#####  
#  
# 1. In this section, you will create your own myLocalSum operation  
#  
#####
```

@localoperation
def myLocalSum(self, locations, subdomains):

Function

A function is defined using the *def* keyword and can accept arguments and execute code statements. Statements within functions must be indented.

```
def function_name(parameters):  
    Python code
```

1b – Define your Function

All PCML functions accept the same parameters:

self – (Python standard) To access Operation-level information

locations – List of locations to process using this function

subdomains – List of subdomains associated with the locations

@localoperation

def myLocalSum(self, locations, subdomains):

Function

A function is defined using the *def* keyword and can accept arguments and execute code statements. Statements within functions must be indented.

```
def function_name(parameters):  
    Python code
```

1c – Create your location variables

```
#####  
#  
# 1. In this section, you will create your own myLocalSum operation  
#  
#####
```

```
@localoperation  
def myLocalSum(self, locations, subdomains):  
    sum = 0
```

Variables

Variables can be assigned values using the equal sign (just like in math class).

1c – Create your location variables

```
#####  
#  
# 1. In this section, you will create your own myLocalSum operation  
#  
#####
```

```
@localoperation  
def myLocalSum(self, locations, subdomains):  
    sum = 0  
    for loc in locations:
```

For loops

This for loop will go over each location (loc) one by one.

Dictionary

```
adictionary = { 'x': 41.1, 'y': -81.3, 'v': 10 }  
  
print adictionary['x']    # Print value for this key  
                           # which would print 41.1
```

Dictionaries

Dictionaries store multiple elements in a single data structure.

Elements consist of a **name**:

'x', 'y', or 'v' and a **value**: 41.1, -81.3, 10.

1d – Add each location value to sum

All PCML locations are dictionaries with 3 keys:

x – X coordinate
y – Y coordinate
v – Value at the location

@localoperation

```
def myLocalSum(self, locations, subdomains):  
    sum = 0  
    for loc in locations:  
        sum = sum + loc['v']
```

Add up each value ('v') from all the locations (loc)

Python Statement Breakdown

sum = sum + loc['v']

Value

Access the location's value using the name 'v'. This is the standard way to access location values in PCML.

Python Statement Breakdown

```
sum = sum + loc['v']
```



Add Sum to Value

The sum variable holds the current sum.

Add the current sum and the location's value together

Python Statement Breakdown

```
sum = sum + loc['v']
```



New Sum Variable

Now sum will be assigned
the result from adding
the current sum and
the location's value

1e – Return your result

```
#####  
#  
# 1. In this section, you will create your own myLocalSum operation  
#  
#####
```

```
@localoperation  
def myLocalSum(self, locations, subdomains):  
    sum = 0  
    for loc in locations:  
        sum = sum + loc['v']  
    return sum
```

Function return

A function can return a value to the code that called it.

Awesomeness (Almost)

You did it!

@localoperation

```
def myLocalSum(self, locations, subdomains):  
    sum = 0  
    for loc in locations:  
        sum = sum + loc['v']  
    return sum
```

myLocalSum – What you wrote

loc['v'] $1 + 2 = 3$ return sum

layera

1	2
3	4

+

layerb

2	4
1	6

=

myresult

3	6
4	10

Your Local Sum Operation

Call your new Operation

```
#####
```

```
#
```

```
# 2. In this section, call your myLocalSum operation
```

```
#   and save the results to a 'myresult' layer
```

```
#
```

```
#####
```

```
myresult=myLocalSum(layera,layerb)
```

**Indent 4 spaces to
Align with comments**

Next, we need to look for bugs

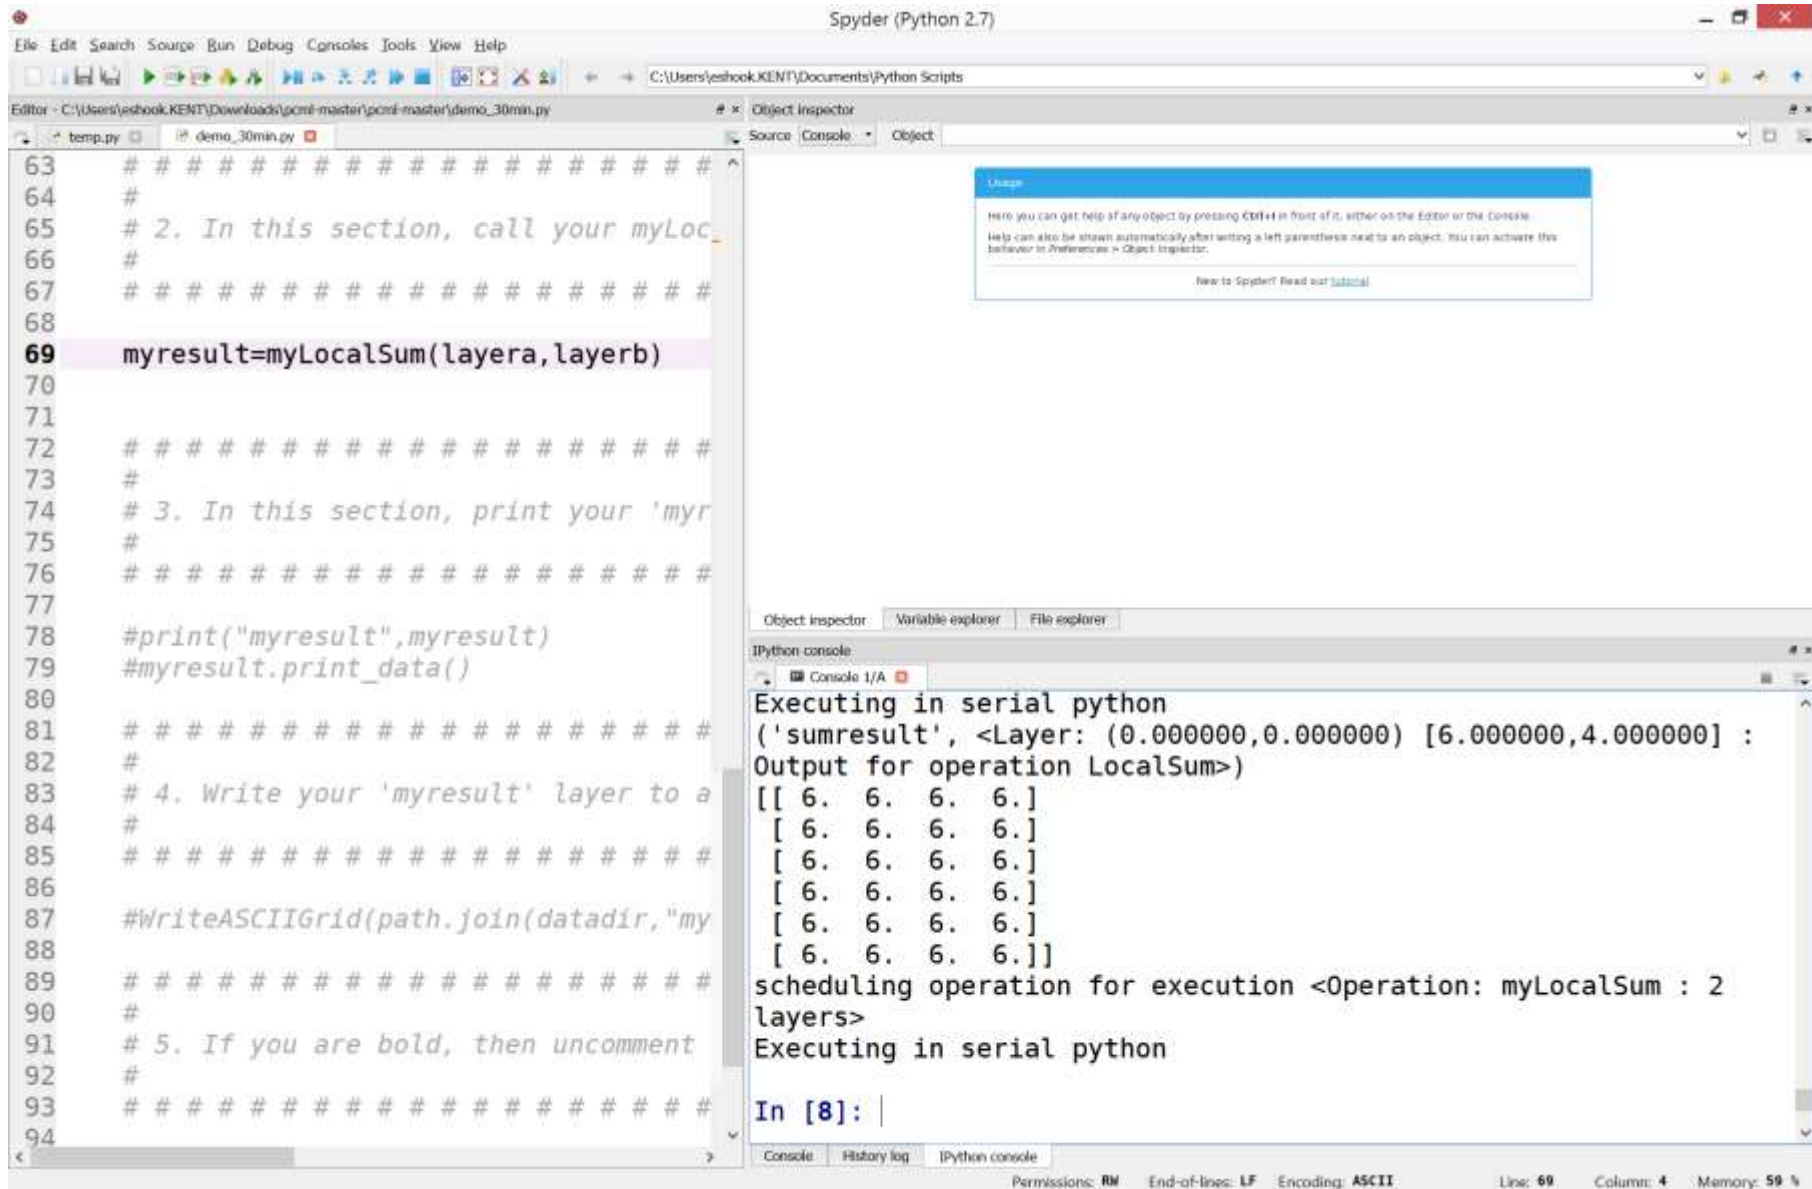


Not these bugs

**These are
bugs and errors
that you may see**

```
$ python demo_30min.py
  File "demo_30min.py", line 4
    ^
SyntaxError: invalid syntax
```

Run the Demo Program



The screenshot shows the Spyder Python IDE interface. The main editor window displays a Python script with the following visible code:

```
63 #####
64 #
65 # 2. In this section, call your myLoc
66 #
67 #####
68
69 myresult=myLocalSum(layera,layerb)
70
71
72 #####
73 #
74 # 3. In this section, print your 'myr
75 #
76 #####
77
78 #print("myresult",myresult)
79 #myresult.print_data()
80
81 #####
82 #
83 # 4. Write your 'myresult' layer to a
84 #
85 #####
86
87 #WriteASCIIGrid(path.join(datadir,"my
88
89 #####
90 #
91 # 5. If you are bold, then uncomment
92 #
93 #####
94
```

The IPython console at the bottom right shows the execution output:

```
Executing in serial python
('sumresult', <Layer: (0.000000,0.000000) [6.000000,4.000000] :
Output for operation LocalSum>)
[[ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]]
scheduling operation for execution <Operation: myLocalSum : 2
layers>
Executing in serial python
In [8]: |
```

The status bar at the bottom indicates: Permissions: RW, End-of-lines: LF, Encoding: ASCII, Line: 69, Column: 4, Memory: 59 %.

Spyder editor shown (any editor will do)

Now Print and Write the Output

```
#####  
# 3. In this section, print your 'myresult' layer  
#  
#####
```

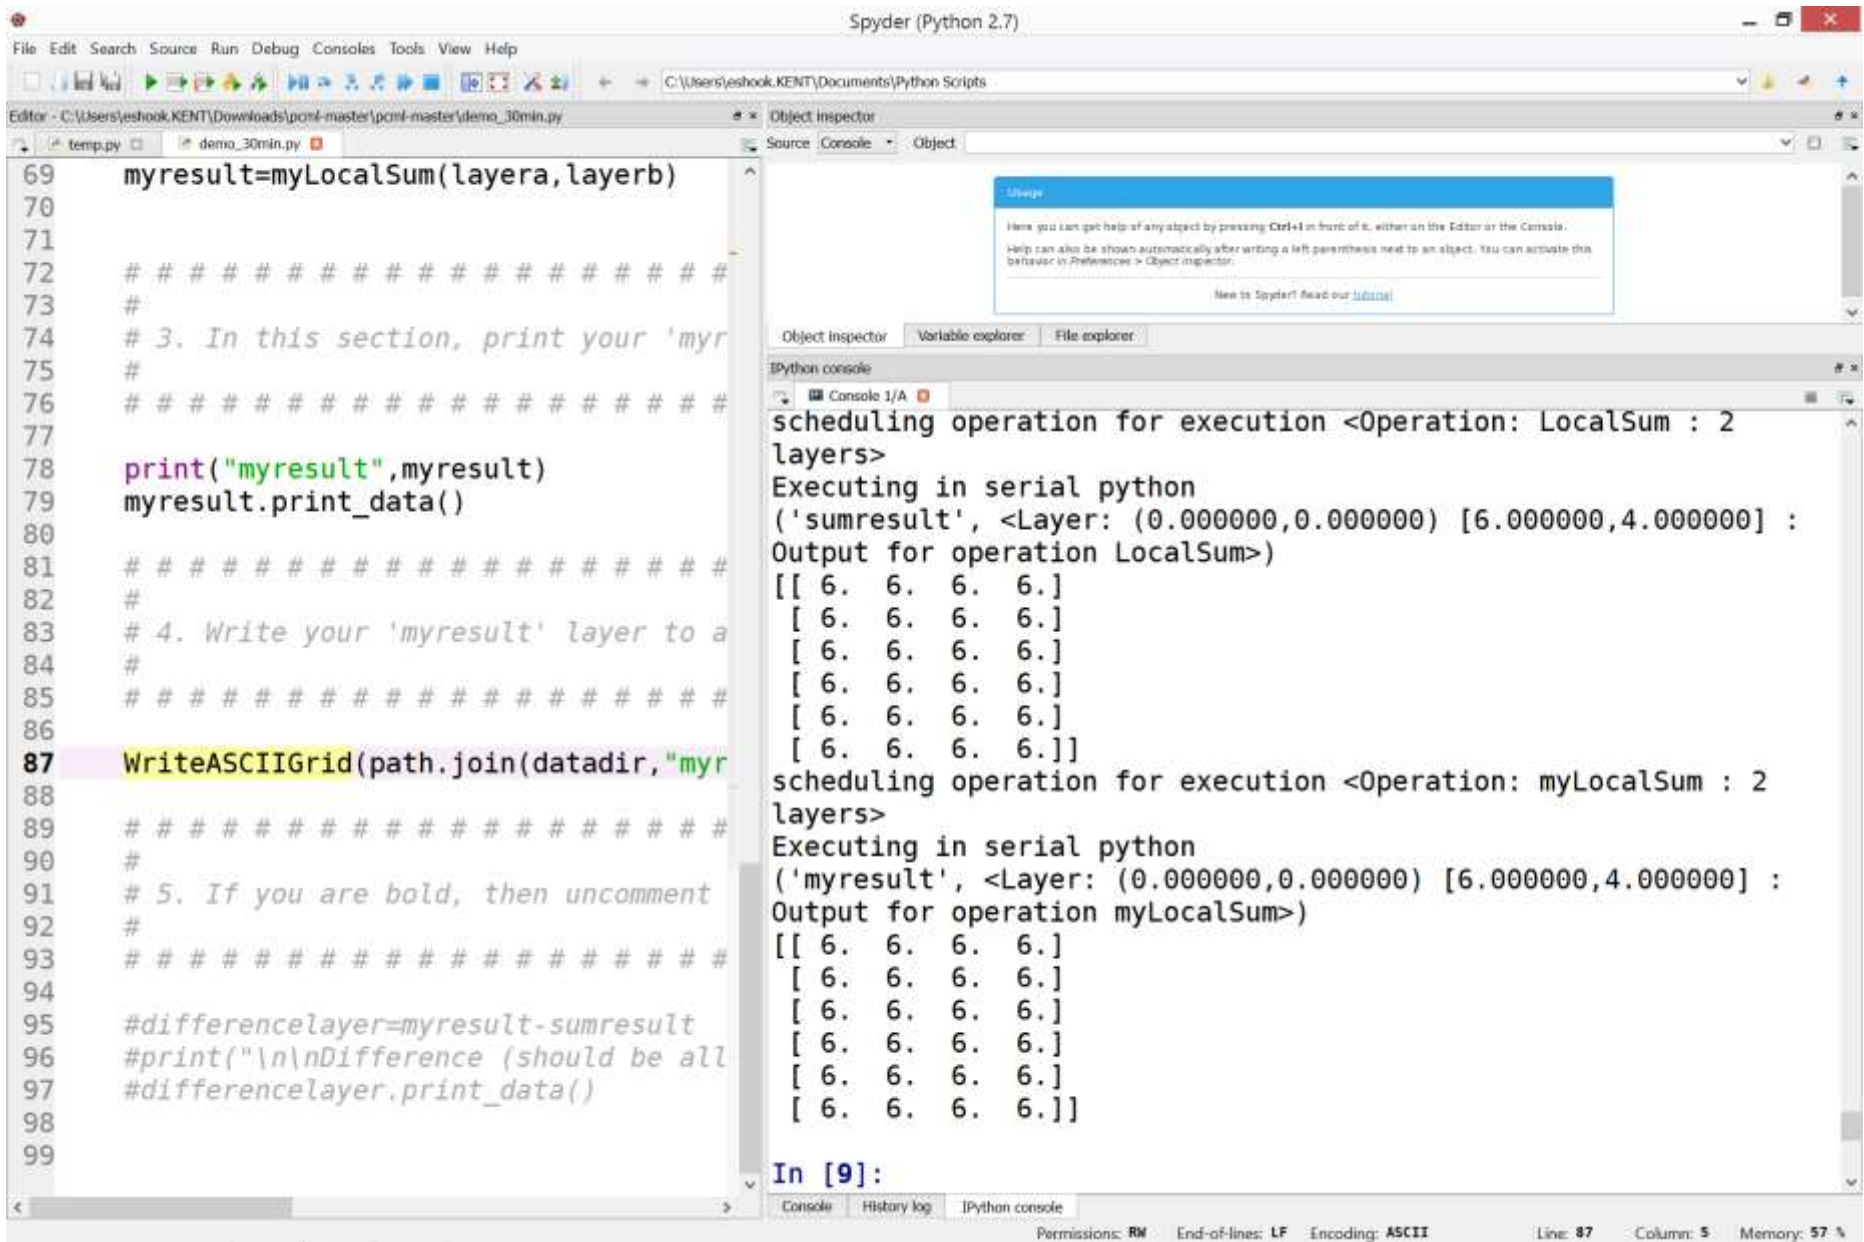
```
print("myresult",myresult)  
myresult.print_data()
```

```
#####  
# 4. Write your 'myresult' layer to a file named 'myresult.asc'  
#  
#####
```

```
WriteASCIIGrid(path.join(datadir,"myresult.asc"), myresult)
```



Now Print and Write the Output



The screenshot displays the Spyder Python IDE interface. The main editor window shows a Python script with the following code:

```
69 myresult=myLocalSum(layera,layerb)
70
71
72 #####
73 #
74 # 3. In this section, print your 'myr
75 #
76 #####
77
78 print("myresult",myresult)
79 myresult.print_data()
80
81 #####
82 #
83 # 4. Write your 'myresult' layer to a
84 #
85 #####
86
87 WriteASCIIGrid(path.join(datadir,"myr
88
89 #####
90 #
91 # 5. If you are bold, then uncomment
92 #
93 #####
94
95 #difference=layera-sumresult
96 #print("\n\nDifference (should be all
97 #difference.print_data()
98
99
```

The IPython console on the right shows the execution output:

```
scheduling operation for execution <Operation: LocalSum : 2
layers>
Executing in serial python
('sumresult', <Layer: (0.000000,0.000000) [6.000000,4.000000] :
Output for operation LocalSum>)
[[ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]]
scheduling operation for execution <Operation: myLocalSum : 2
layers>
Executing in serial python
('myresult', <Layer: (0.000000,0.000000) [6.000000,4.000000] :
Output for operation myLocalSum>)
[[ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]]
In [9]:
```

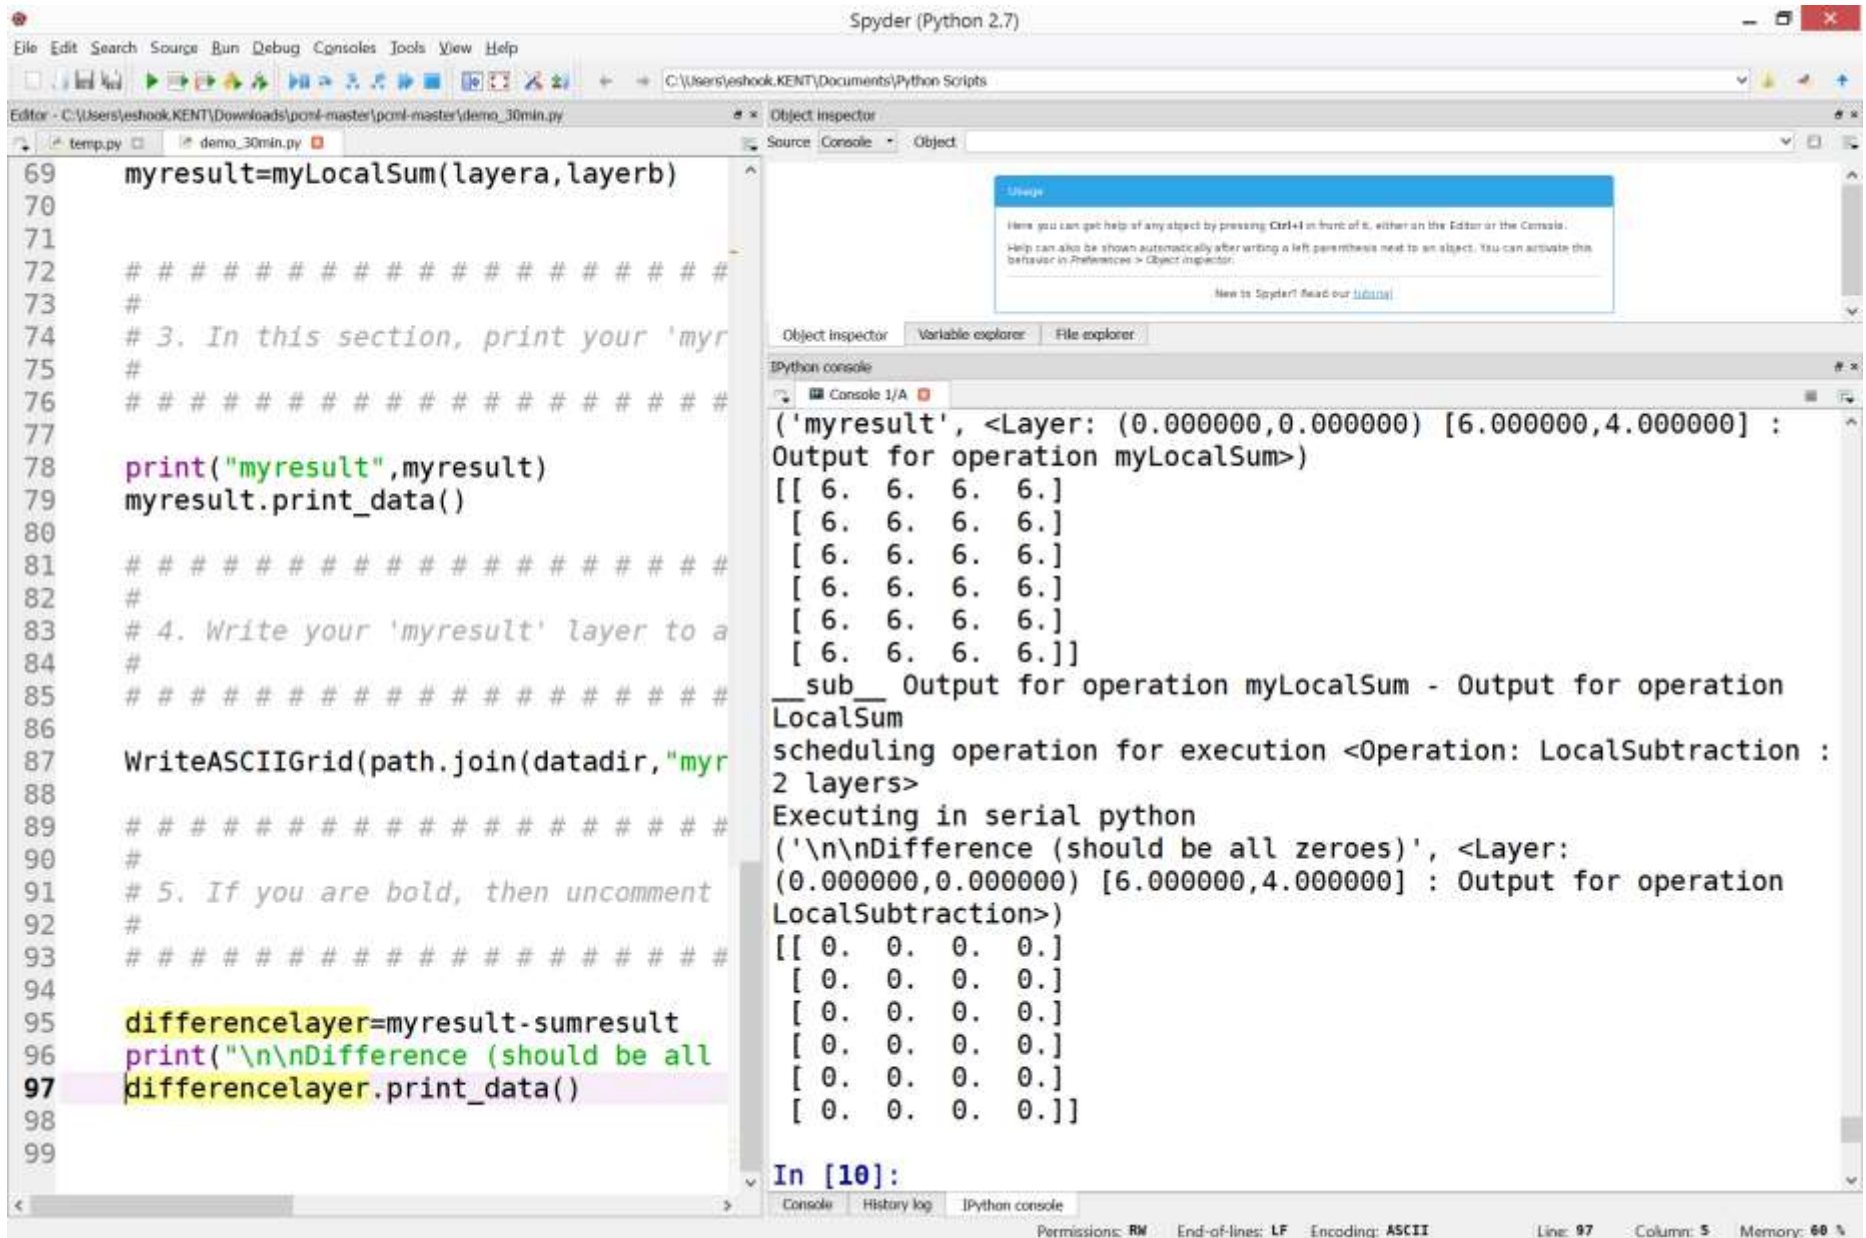
The status bar at the bottom indicates: Permissions: RW, End-of-lines: LF, Encoding: ASCII, Line: 87, Column: 5, Memory: 57 %.

The Ultimate Test: The Difference

```
#####  
#  
# 5. If you are bold, then uncomment the next three lines of code  
#   to see if the difference is zero!  
#  
#####  
  
differencelayer=myresult-sumresult  
print("\n\nDifference (should be all zeroes)",differencelayer)  
differencelayer.print_data()
```



Success!



```
69 myresult=myLocalSum(layera,layerb)
70
71
72 #####
73 #
74 # 3. In this section, print your 'myr
75 #
76 #####
77
78 print("myresult",myresult)
79 myresult.print_data()
80
81 #####
82 #
83 # 4. Write your 'myresult' layer to a
84 #
85 #####
86
87 WriteASCIIGrid(path.join(datadir,"myr
88
89 #####
90 #
91 # 5. If you are bold, then uncomment
92 #
93 #####
94
95 differencelayer=myresult-sumresult
96 print("\n\nDifference (should be all
97 differencelayer.print_data()
98
99
```

Object Inspector

Variable explorer

File explorer

IPython console

Console 1/A

```
('myresult', <Layer: (0.000000,0.000000) [6.000000,4.000000] :
Output for operation myLocalSum>)
[[ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]
 [ 6.  6.  6.  6.]]
__sub__ Output for operation myLocalSum - Output for operation
LocalSum
scheduling operation for execution <Operation: LocalSubtraction :
2 layers>
Executing in serial python
('\n\nDifference (should be all zeroes)', <Layer:
(0.000000,0.000000) [6.000000,4.000000] : Output for operation
LocalSubtraction>)
[[ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]]

In [10]:
```

Permissions: RW End-of-lines: LF Encoding: ASCII Line: 97 Column: 5 Memory: 60 %

Other Things to Try

Change the Iteration strategy *

`LocalSum(layera,layerb,iteration=columnmajoriteration)`

Change the Decomposition strategy *

`LocalSum(layera,layerb,decomposition=columndecomposition)`

Change the Executor Type from Parallel to Serial

`PCMLConfig.exectype = ExecutorType.serialpython`

Change the number of cores

`PCMLConfig.num_procs = 1`

Use a different operation

Take a look at the library of operations (in **pcml/lib/**)

* add **print(loc)** in the for loop to see the changes in Iteration and Decomposition



Welcome to the (awesome) group of PCML developers!

- * PCML is an open-source project
and we welcome your contributions



Acknowledgements

A huge thanks to (the growing number of) PCML contributors:

Michael Hodgson

Shaowen Wang

Babak Behzad

Kiumars Soltani

April Hiscox

Zhengliang Feng

Emil Shirima

Jayakrishnan Ajayakumar

Sandeep Vutla

Gowtham Kukkadapu

Suman Jindam

This work was in-part funded by the CyberGIS Fellows program

This work used the Kent State University Arts and Sciences Cluster



Thank you!

Questions?

Please contact
Eric Shook
eshook@gmail.com